

Yuan Huang 黄远

huangyuan4629@gmail.com

Homepage • Google Scholar • GitHub

Northeastern University, Shenyang, China

RESEARCH INTERESTS

Natural language processing, with a focus on **measurement validity** — whether the benchmarks and automatic judges used to rank models measure what they claim to. One result recurs across my projects: **surface competence $\neq$ underlying competence**. I want to measure that gap, explain it mechanistically, and close it.

EDUCATION

Northeastern University, China — B.Eng., Computer Science and Technology 2023 – 2027 (expected)

- CGPA **4.05 / 5.00** (90 / 100) · **Rank 5 / 100** in major cohort
- Advised by Shiqi Zhao (Assistant Professor, Northeastern University) since 2024

PUBLICATIONS

[1] Beyond Survival: Evaluating LLMs in Social Deduction Games with Human-Aligned Strategies

Zirui Song*, **Yuan Huang***, Junchang Liu*, et al. (* *equal contribution*)

EMNLP 2026, Main Conference • co-first author • arXiv:2510.11389 • Dataset

[2] Geolocation with Real Human Gameplay Data: A Large-Scale Dataset and Human-Like Reasoning Framework

Zirui Song, Jingpu Yang, **Yuan Huang**, et al.

Preprint • arXiv:2502.13759 • Dataset • Code

RESEARCH EXPERIENCE

EditJudgeBias — counterfactual audit of multimodal image-editing judges

2026 – present

Sole first author • in preparation, targeting ICLR 2027

- **Design.** Admit only perturbations that can be shown not to have changed edit quality: preservation verification is an admission gate, not a post-hoc check, and every cue carries a zero-dose sham control, so null results are readable.
- **Scale.** 13 cue conditions across four injection sites, 5 published judges plus an independent replication arm — **97,973 judgments**, 0.300% parse failure.
- **Findings.** Reordering the two candidates alone flips decisions at **38.2 $\times$ the judge’s self-noise**; of three mitigation layers only the protocol-level one works, and it works by **abstention, not debiasing**.
- **Rigor.** 1,307 unit tests and 109 sha256-stamped frozen artifacts; the audit overturned four of my own conclusions.

WereBench / WereAlign — human-aligned evaluation of LLM social reasoning

EMNLP 2026 Main

Co-first author • inception through submission

- Benchmark built from **100+ hours** of professional human social-deduction gameplay (**32.4M utterance tokens**, 30 roles, 15 rule variants), scored against **winning-side MVP trajectories** rather than LLM self-play; 5 verbal dimensions and 2 decision tasks.
- **Half of the evaluated models score below 0.50** (best 0.720). Deception reasoning and counterfactual trade-off are weakest while persuasive generation stays strong — **fluency $\neq$ strategic correctness**.

GeoComp / GeoCoT / GeoEval — geolocation from real human gameplay data

Preprint

Co-author (3 of 9)

- Dataset from a public geo-guessing platform — **740K users, 2.7M locations, 25M human response records** — with human accuracy calibrating item difficulty. Multi-step geographic reasoning improved standard metrics by up to **25%**; scoring the reasoning process rather than only the final answer added a further **9%**.

HONORS & AWARDS

- **China National Scholarship** — national level · 2023–2024
- Northeastern University Comprehensive Scholarship · 2024–2026
- **Third Prize, National Finals**, Lanqiao Cup Programming Competition · 2024

SKILLS

- Python, C++, Java · PyTorch, Hugging Face Transformers · backend engineering intern, 2024 (Java / Spring Boot)
- IELTS Academic — taken 4 September 2026 (score to follow) · CET-6 534