

Yuan Huang 黄远

huangyuan4629@gmail.com · +86 191 6050 6039
Homepage · Google Scholar · GitHub

Northeastern University, Shenyang, China

INTERESTS

Trustworthy and human-centered evaluation of large (multimodal) language models. Specifically: benchmark design grounded in large-scale real human behavior rather than model self-play; reliability and bias of LLM-as-a-judge; **mechanistic accounts of why a model's judgment fails**; and social and strategic reasoning in LLM agents.

EDUCATION

B.Eng. in Computer Science and Technology — Northeastern University, China

2023 – 2027 (expected mid-2027) · Shenyang, China

— CGPA **4.05 / 5.00** (90 / 100) · **Rank 5 / 100** in major cohort

PUBLICATIONS

* denotes equal contribution.

Beyond Survival: Evaluating LLMs in Social Deduction Games with Human-Aligned Strategies

Zirui Song*, Yuan Huang*, Junchang Liu*, et al.

EMNLP 2026 MAIN · arXiv:2510.11389 · co-first author

Geolocation with Real Human Gameplay Data: A Large-Scale Dataset and Human-Like Reasoning Framework

Zirui Song, Jingpu Yang, Yuan Huang, et al.

PREPRINT · arXiv:2502.13759

Inside the Unfair Judge: A Mechanistic Interpretability Account of MLLM-as-Judge Bias

Yuan Huang — sole first author

IN PREPARATION · targeting ICLR 2027 (abstract deadline

18 Sep 2026)

RESEARCH

Mechanistic Interpretability of MLLM-as-Judge Bias — independent first-author project

2026 – present · Sole first author: conception, experimental design, implementation, analysis, writing

- Showed that multimodal LLM judges are steered by **non-visual** factors while the image stays **pixel-identical**: attaching a caption line lowers the score; inserting a social cue (“most reviewers preferred it”) into the prompt raises it.
- Gave a **mechanistic account of how the bias forms inside the model**, rather than reporting the behavioral effect alone.
- Evaluated **three widely used debiasing strategies**; **none** restored the correct judgment.

WereBench / WereAlign — human-aligned evaluation of LLM social reasoning

EMNLP 2026 Main · Co-first author; from project inception through submission

- Built a benchmark from **100+ hours** of professionally produced human social-deduction gameplay (**32.4M utterance tokens**, 30 roles, 15 rule variants, multimodal), using **winning-side MVP strategy trajectories as ground truth** — in contrast to prior LLM self-play data.
- Designed a 5-dimension verbal evaluation and 2 decision tasks; contributed to research design, experiments, evaluation, and iterative refinement throughout.
- **Half of the evaluated SOTA models score below 0.50**; the best reaches only **0.720**. Weakest dimensions are **deception reasoning** and **counterfactual trade-off**, while persuasive generation stays strong — i.e. **fluency ≠ strategic correctness**.
- Two controlled interventions separated distinct failure paths, showing that models can mistake another speaker’s persuasive utterance for an instruction.

GeoComp / GeoCoT / GeoEval — geolocation from real human gameplay data

Preprint · Co-author (3 of 9)

- Contributed to a dataset from a public geo-guessing platform: **740K users, 2.7M locations, 25M human response records**, with human accuracy used to calibrate item difficulty.
- Multi-step geographic reasoning improved standard geolocation metrics by up to **25%**; scoring the **reasoning process** rather than only the final answer added a further **9%**.

HONORS

China National Scholarship — national level, top ~1% nationally · 2023–2024

Northeastern University Comprehensive Scholarship · 2023–2026

Third Prize, National Finals, Lanqiao Cup Programming Competition · 2024

SKILLS

LANGUAGES

C++, Python, Java

ML / RESEARCH

PyTorch, Hugging Face Transformers

HUMAN LANGUAGES

Chinese (native) · English — CET-6 534; IELTS Academic taken 4 September 2026

ALSO

Backend engineering intern, 2024 — Java / Spring Boot.